

Informed Search (or Heuristic Search)

Vincent A. Cicirello, Ph.D.

Professor of Computer Science

cicirelv@stockton.edu

<https://www.cicirello.org/>



1

Lesson 1

INFORMED VS UNINFORMED SEARCH



2

Overview

- What is the difference between *informed* search and *uninformed* search?
- Heuristics: what are they? And what are they good for?
- Best-first greedy search
- A* Search
 - A* is complete
 - Admissible heuristics
 - A*'s dark side
- Iterative Deepening A* (IDA*) Search



3

Uninformed Search

- Uninformed search algorithms are search algorithms that don't rely on any problem-dependent knowledge.
 - They are *uninformed* of the problem they are solving.
 - All of the search algorithms that we've seen so far: BFS, DFS, Least-Cost BFS, UCS, Depth-Limited DFS, Iterative Deepening.
- All of these have a worst-case runtime that is exponential in the average branching factor of the problem.
 - Not much we can do about that in general.
 - The problems of interest to A.I. are often NP-Complete or NP-Hard.



4

Informed Search

- Informed search algorithms use problem-dependent knowledge, in the form of a heuristic, for guidance.
 - They are informed of the problem they are solving.
 - A heuristic is used to attempt to more intelligently choose which successors appear to be more promising than the others.
- The informed search algorithms that we'll study also have worst-case runtime that is exponential in problem's average branching factor.
 - Not much we can do about that in general.
 - The problems of interest to A.I. are often NP-Complete or NP-Hard.
- With good heuristics, they tend to avoid exponentially-long runs.



5

Heuristics

Suppose in addition to the standard search specification we also have a *heuristic*.

A heuristic function maps a state onto an estimate of the cost to the goal from that state.

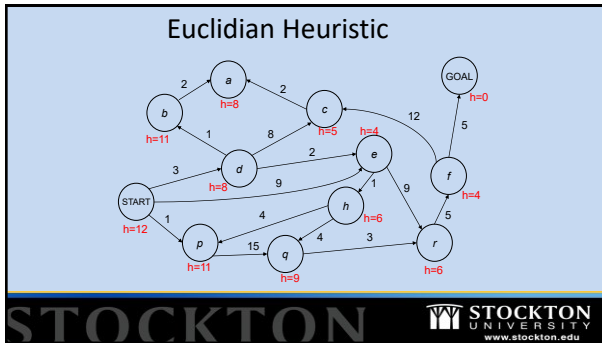
Can you think of examples of heuristics?

- E.G. for the 8-puzzle?
- E.G. for planning a path through a maze?

Denote the heuristic by a function $h(s)$ from states to a cost value.



6



7

Eight Puzzle Heuristic

- A few example heuristics for the 8-puzzle:
 - Number of tiles in the wrong spots
 - Sum of Manhattan distances
- Is one of these better than the other?
 - Yes: Sum of Manhattan distances

Example State

1	5	
2	6	3
7	4	8

Goal State

1	2	3
4	5	6
7	8	

STOCKTON UNIVERSITY
www.stockton.edu

8

Maze Solving

- A few heuristics for maze solving:
 - Euclidian distance
 - Manhattan distance
- Is one of these better than the other?
 - Yes: Manhattan distance.

STOCKTON UNIVERSITY
www.stockton.edu

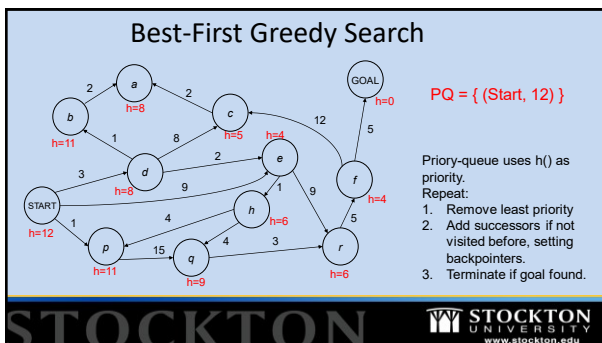
9

Lesson 2

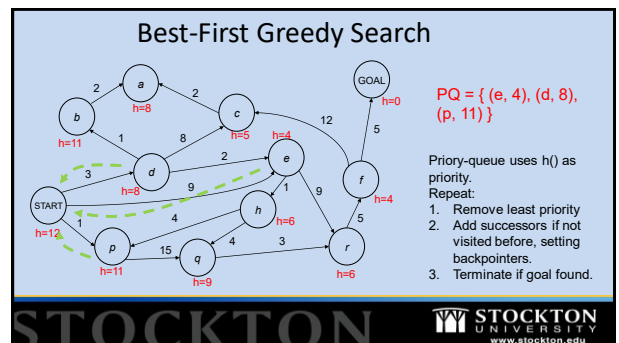
BEST-FIRST GREEDY SEARCH

STOCKTON UNIVERSITY
www.stockton.edu

10



11



12

Best-First Greedy Search

PQ = { (r, 6), (h, 6), (d, 8), (p, 11) }

Priority-queue uses $h()$ as priority.
Repeat:
1. Remove least priority
2. Add successors if not visited before, setting backpointers.
3. Terminate if goal found.

STOCKTON UNIVERSITY
www.stockton.edu

13

Best-First Greedy Search

PQ = { (f, 4), (h, 6), (d, 8), (p, 11) }

Priority-queue uses $h()$ as priority.
Repeat:
1. Remove least priority
2. Add successors if not visited before, setting backpointers.
3. Terminate if goal found.

STOCKTON UNIVERSITY
www.stockton.edu

14

Best-First Greedy Search

PQ = { (Goal, 0), (c, 5), (h, 6), (d, 8), (p, 11) }

Priority-queue uses $h()$ as priority.
Repeat:
1. Remove least priority
2. Add successors if not visited before, setting backpointers.
3. Terminate if goal found.

STOCKTON UNIVERSITY
www.stockton.edu

15

Best-First Greedy Search

PQ = { (c, 5), (h, 6), (d, 8), (p, 11) }

Goal is removed from PQ, so search returns the path implied by backpointers:
Start, e, r, f, Goal
With a cost of $9+9+5+5=28$

Actual optimal path is:
Start, d, e, h, q, r, f, Goal
Cost: $3+2+1+4+3+5+5=23$

STOCKTON UNIVERSITY
www.stockton.edu

16

Best First "Greedy" Search

```

Init-PriQueue(PQ)
Insert-PriQueue(PQ, START, h(START))
while (PQ is not empty and PQ does not contain a goal state)
  (s, h) := Pop-least(PQ)
  foreach s' in succs(s)
    if s' is not already in PQ and s' never previously been visited
      Insert-PriQueue(PQ, s', h(s'))
  
```

Algorithm	Complete	Optimal	Time	Space
Best First Greedy Search				

STOCKTON UNIVERSITY
www.stockton.edu

17

Best First "Greedy" Search

```

Init-PriQueue(PQ)
Insert-PriQueue(PQ, START, h(START))
while (PQ is not empty and PQ does not contain a goal state)
  (s, h) := Pop-least(PQ)
  foreach s' in succs(s)
    if s' is not already in PQ and s' never previously been visited
      Insert-PriQueue(PQ, s', h(s'))
  
```

Algorithm	Complete	Optimal	Time	Space
Best First Greedy Search	Y			

STOCKTON UNIVERSITY
www.stockton.edu

18

Best First “Greedy” Search

```

Init-PriQueue(PQ)
Insert-PriQueue(PQ, START, h(START))
while (PQ is not empty and PQ does not contain a goal state)
  (s, h) := Pop-least(PQ)
  foreach s' in succs(s)
    if s' is not already in PQ and s' never previously been visited
      Insert-PriQueue(PQ, s', h(s'))
    
```

Algorithm	Complete	Optimal	Time	Space
Best First Greedy Search	Y	N		

19

Best First “Greedy” Search

```

Init-PriQueue(PQ)
Insert-PriQueue(PQ, START, h(START))
while (PQ is not empty and PQ does not contain a goal state)
  (s, h) := Pop-least(PQ)
  foreach s' in succs(s)
    if s' is not already in PQ and s' never previously been visited
      Insert-PriQueue(PQ, s', h(s'))
    
```

Algorithm	Complete	Optimal	Time	Space
Best First Greedy Search	Y	N	$O(\log(Q) * \min(BN, B^{LMAX}))$	$O(\min(N, B^{LMAX}))$

20


Lesson 3

A* SEARCH

21


Let's Make “Best first Greedy” Look Stupid!

- $PQ = \{(S, 4)\}$
- Expand S
 - $PQ = \{(A, 3)\}$
- Expand A
 - $PQ = \{(C, 1), (B, 2)\}$
- Expand C
 - $PQ = \{(G, 0), (B, 2)\}$
- Expand G... found the goal... solution path S,A,C,G has cost 94...
- But a shorter path exists... S,A,B,C,G with cost 6



22

Let's Make “Best first Greedy” Look Stupid!



- Best-first greedy is clearly not guaranteed to find optimal
- Question: What can we do to avoid the stupid mistake?

23

A* - The Basic Idea

- **Best-first greedy:** When you expand a node s , take each successor s' and place it on PriQueue with priority $h(s')$
- **A*:** When you expand a node s , take each successor s' and place it on PriQueue with priority

$$(\text{Cost of getting to } s) + h(s) \quad (1)$$

$$\text{Let } g(s) = \text{Cost of getting to } s \quad (2)$$

and then define...

$$f(s) = g(s) + h(s) \quad (3)$$

24


A* Looking Non-Stupid

When we're at state A...

We have 2 successors, B and C...
 $f(B) = g(B) + h(B) = 3 + 2 = 5$
 $f(C) = g(C) + h(C) = 92 + 1 = 93$
 $PQ = \{ (B, 5), (C, 93) \}$

When we're at state B...

One successor, C...
 $f(C) = g(C) + h(C) = 4 + 1 = 5$
 $PQ = \{ (C, 5) \}$




25

When should A* terminate?

Idea: As soon as it generates a goal state?

Look at this example:




26

When should A* terminate?

Idea: As soon as it generates a goal state?

Look at this example:




- Expand S
 $PQ = \{ (B, 4), (A, 8) \}$
- Expand B
 $PQ = \{ (C, 4), (A, 8) \}$
- Expand C
 $PQ = \{ (D, 4), (A, 8) \}$
- Expand D
 $PQ = \{ (A, 8), (G, 10) \}$
- A goal state has been generated

27

Correct A* termination rule:


A* Terminates Only When a Goal State Is Removed from the Priority Queue



28

Correct A* termination rule:

A* Terminates Only When a Goal State Is Removed from the Priority Queue




- Expand S
 $PQ = \{ (B, 4), (A, 8) \}$
- Expand B
 $PQ = \{ (C, 4), (A, 8) \}$
- Expand C
 $PQ = \{ (D, 4), (A, 8) \}$
- Expand D
 $PQ = \{ (A, 8), (G, 10) \}$
- Expand A
 $PQ = \{ (G, 8) \}$
- Expand G
 • Found shorter path

29

A* revisiting states

Another question: What if A* revisits a state that was already expanded, and discovers a shorter path?



In this example a state that had been expanded gets re-expanded.

- Expand S
 $PQ = \{ (B, 4), (A, 8) \}$
- Expand B
 $PQ = \{ (C, 4), (A, 8) \}$
- Expand C
 $PQ = \{ (D, 4), (A, 8) \}$
- Expand D
 $PQ = \{ (A, 8), (G, 10) \}$
- Expand A
 $PQ = \{ (C, 3.5), (G, 10) \}$
 - Found less costly path to C.
- Expand C
 $PQ = \{ (D, 3.5), (G, 10) \}$
- Expand D
 $PQ = \{ (G, 9.5) \}$
- Expand G
 • Done
 • Path: S, A, C, D, G

30

A* revisiting states

Another question: What if A* visits a state that is already on the queue?

- Expand S
 - PQ = { (B,4), (A,8) }
- Expand B
 - PQ = { (A,8), (C,10) }
- Expand A
 - PQ = { (C, 9.5) }
 - Found less costly path to C.
- Expand C
 - PQ = { (D,3.5) }
- Expand D
 - PQ = { (G, 9.5) }
- Expand G
 - Done
 - Path: S, A, C, D, G

note that this h value has changed from previous slide.

31

A* Search

- Priority queue PQ begins empty.
- V (= set of previously visited (state,f,backpointer)-triples) begins empty.
- Put S into PQ and V with priority $f(S) = g(S) + h(S)$
 - $h(S)$ because $g(start) = 0$
- Is PQ empty?
 - Yes? There's no solution
 - No? Remove node with lowest $f(s)$ from queue. Call it s.
 - If s is a goal, stop and report success.
 - "expand" s : For each s' in $successors(s)$...
 - Let $f' = g(s') + h(s') = g(s) + cost(s,s') + h(s')$
 - If s' not seen before, or s' previously expanded with $f(s') > f'$, or s' currently in PQ with $f(s') > f'$
 - Then Place/promote s' on priority queue with priority f' and update V to include (state= s' , f' , BackPtr=s).
 - Else Ignore s'

Reminder: $g(s)$ is cost of shortest known path to s

Reminder: $h(s)$ is a heuristic estimate of cost to a goal from s

32

Is A* Guaranteed to Find the Optimal Path?

- Expand S
 - PQ = { (G, 3), (A, 7) }
- Expand G
 - Done
 - Found path: S, G
 - Cost 3
- But a less costly path exists
 - S, A, G
 - Cost 2

Nope. And this example shows why not.

Next, we'll look at what we can do guarantee an optimal solution. We'll need to put requirements on the heuristic.

33

Lesson 4

ADMISSIBLE HEURISTICS

34

A* is not Optimal (at least not in general)

- Expand S
 - PQ = { (G, 3), (A, 7) }
- Expand G
 - Done
 - Found path: S, G
 - Cost 3
- But a less costly path exists
 - S, A, G
 - Cost 2

Nope. And this example shows why not.

If we put some requirements on the heuristic, then we can get A* to find the optimal solution.

35

Admissible Heuristics

- Write $h^*(s)$ = the true minimal cost to goal from s.
- A heuristic $h(s)$ is **admissible** if $h(s) \leq h^*(s)$ for all states s.
- An admissible heuristic is guaranteed never to overestimate cost to goal.
- An admissible heuristic is optimistic.

36

8-Puzzle Example

Example State

1	5	
2	6	3
7	4	8

Goal State

1	2	3
4	5	6
7	8	

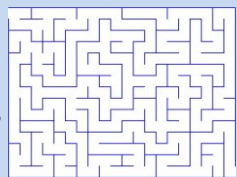
Which of the following are admissible heuristics?

- $h(s)$ = Number of tiles in wrong position in state s
- $h(s) = 0$
- $h(s)$ = Sum of Manhattan distances between each tile and its goal location
- $h(s) = 1$
- $h(s) = \min(2, h^*[s])$
- $h(s) = h^*(n)$
- $h(s) = \max(2, h^*[s])$

37

Maze Solving

- A few **admissible** heuristics for maze solving:
 - Euclidian distance
 - Manhattan distance
- Both are admissible.
 - Without obstacles, shortest path between two points is a straight line (Euclidian distance), so will clearly never overestimate.
 - Only able to move left-right-up-down, so without walls, an optimal path would be to walk left-right straight line to x of goal, and then up-down to y of goal.



38

A* and Optimality

- A* is complete.
 - Guaranteed to find a solution if one exists and to terminate otherwise.
- A* with an admissible heuristic is optimal.
 - Guaranteed to find the optimal path from the start to the goal.

39

Lesson 5

DOMINATING HEURISTICS

40

Comparing A* with Iterative Deepening

From Russell and Norvig

For 8-puzzle, average number of states expanded over 100 randomly chosen problems in which optimal path is length...

	...4 steps	...8 steps	...12 steps
Iterative Deepening (see slides on uninformed search)	112	6,300	3.6×10^6
A* search using "number of misplaced tiles" as the heuristic	13	39	227
A* using "Sum of Manhattan distances" as the heuristic	12	25	73

41

Comments

- At first sight might look like even "number of misplaced tiles" is a great heuristic. But probably $h(\text{state})=0$ would also do much much better than ID, so the difference is mainly to do with ID's big problem of expanding the same state many times, not the use of a heuristic.
- Judging solely by "number of states expanded" does not account for overhead of maintaining hash tables and priority queue for A*, though it's pretty clear here that this won't dramatically change the results.

There are only a couple hundred thousand states for the entire eight puzzle

over 100 in problems in which path is length...

	steps	...12 steps
A* search using "number of misplaced tiles" as the heuristic	39	227
A* using "Sum of Manhattan distances" as the heuristic	12	73

42

Dominating Heuristics

- Given 2 admissible heuristics for a problem, one might ask whether one is always better?
- Example:** *Is Manhattan distance always better than number of misplaced tiles for the 8-puzzle?* **Yes**
- A heuristic h_2 **dominates** h_1 if for all states s , $h_2(s) \geq h_1(s)$
- Using A^* with heuristic h_2 will never expand more nodes than A^* with h_1 .
- It is always better to use a heuristic function with higher values, provided:
 - it is admissible (never overestimates)
 - and that the computation time for the heuristic is not too large

43

Additional Examples of Dominating Heuristics

- For the eight puzzle (and sliding tile puzzles in general), Manhattan distance dominates number of misplaced tiles.
- What about maze solving (assuming the maze is a grid with horizontal and vertical movement only)?
- Manhattan distance dominates Euclidean distance.
 - Euclidean distance (single straight line cutting through walls, etc) is shortest path if no obstacles and if free to move in any direction
 - Manhattan distance (two straight lines, cutting through walls, etc, one horizontally, and one vertically)

44

Lesson 6

ITERATIVE DEEPENING A^* SEARCH (IDA*)

45

A^* : The Dark Side

- A^* can use lots of memory.
In principle:
 $O(\text{number of states})$
- For really big search spaces, A^* will run out of memory.



46

IDA* : Memory Bounded Search

- Iterative deepening A^* . Actually, pretty different from A^* . Assume costs integer.
 - Do path-checking DFS, not expanding any state with $f(s) > 0$. Did we find a goal? If so, stop.
 - Do path-checking DFS, not expanding any state with $f(s) > 1$. Did we find a goal? If so, stop.
 - Do path-checking DFS, not expanding any state with $f(s) > 2$. Did we find a goal? If so, stop.
 - Do path-checking DFS, not expanding any state with $f(s) > 3$. Did we find a goal? If so, stop.

...keep doing this, increasing the $f(s)$ threshold by 1 each time, until we stop.
...need to also take care that we detect if problem not solvable (e.g., check if we cut-off search with limit)
- This is
 - Complete
 - Guaranteed to find optimal (if heuristic is admissible)
 - More costly than A^* in general.
 - But only requires memory that is linear in path length.**

47