

Local Search: Hill Climbing, Simulated Annealing, and Tabu Search

Vincent A. Cicirello, Ph.D.
Professor of Computer Science

cicirelv@stockton.edu

<https://www.cicirello.org/>



1

Lesson 1

SEARCH FOR OPTIMIZATION PROBLEMS



2

Optimization Problems

- Optimization Problem:
 - A finite, discrete set, S .
 - Some examples:
 - set of all permutations of N elements,
 - set of all graphs with N nodes,
 - set of all set partitions of a set of N elements,
 - set of all possible variable assignments to a set of N variables, etc
 - A value function, $V : S \rightarrow \mathbb{R}$
 - Search Problem (2 variations):
 1. Find the $s \in S$ such that $s = \underset{S'}{\operatorname{argmax}} V(s')$
 2. Find the $s \in S$ such that $s = \underset{S'}{\operatorname{argmin}} V(s')$



3

An example: Boolean Satisfiability

- You have a set of boolean variables (i.e., variables with domain $\{\text{true}, \text{false}\}$).
- You have a boolean expression in Conjunctive Normal Form, which may or may not be satisfiable.
- Problem: Find an assignment of truth values to variables to maximize number of satisfied clauses.

$$(A \vee B \vee C) \wedge (B \vee C \vee \neg D) \wedge (D \vee E \vee F) \wedge (\neg D \vee \neg E \vee \neg F) \wedge$$

$$(\neg A \vee \neg C \vee \neg F) \wedge (E \vee F \vee \neg G) \wedge (\neg A \vee G \vee \neg H) \wedge$$

$$(D \vee \neg E \vee H) \wedge (B \vee I \vee \neg J) \wedge (\neg H \vee I \vee J) \wedge (G \vee \neg I \vee \neg J)$$



4

Constraint Satisfaction vs Constraint Optimization

- Wait! Isn't Boolean Satisfiability a constraint satisfaction problem?
- Shouldn't we be using the search techniques we saw for constraint satisfaction?
- **Constraint Satisfaction:**
 - Find an assignment of values to variables that satisfy all of the constraints.
- **Constraint Optimization:**
 - Find the assignment of values to variables that satisfy as many constraints as you can.
 - The DFS search techniques we saw previously will be too expensive.



5

Constraint Satisfaction vs Constraint Optimization

- Over-constrained Problems:
 - An over-constrained problem is a constraint satisfaction problem in which no solution exists.
 - In other words, it is impossible to satisfy all constraints simultaneously
- Constraint Satisfaction search algorithms in this case will:
 1. Exhaustively search (via DFS) all possible value assignments (even with all of our tricks, constraint propagation, variable/value ordering heuristics), and then
 2. Simply indicate "no solution."



6

Constraint Satisfaction vs Constraint Optimization


- Often need better than "no solution" as result for over-constrained problem.
- Example:
 - Scheduling courses to classrooms.
 - Each class requires certain room characteristics (e.g., computer lab, chalk board vs whiteboards, etc).
 - Each class has other constraints (number of students, time of day, days of the week).
 - Very near room capacity on Stockton's Galloway campus.
 - Room scheduler used by registrar always fails to find rooms for some classes (over-constrained problem).
 - "No solution" would be an unacceptable outcome.
 - It allocates rooms to as many courses as possible, leaving people to figure out the rest (e.g., changing times, changing constraints, etc)

7

Another example: Traveling salesperson

- Traveling salesperson problem (TSP):

- Set of cities.
- Problem: Find the "tour" of the cities with minimum overall length.
- A "tour" is a simple cycle consisting of **all** of the cities.
- A simple cycle begins and ends at the same city, but otherwise contains no duplicates.



8

Traveling Salesperson Problem

- May sound like a purely theoretical problem with a very old-fashioned application (e.g. salespeople don't usually travel around selling anymore).
- TSP has many important real-world applications.
- Package delivery is an example:
 - Instead of cities, you have delivery addresses within a portion of a city, and the address of the truck depot.
 - You want the delivery driver to visit each address exactly once, returning to the depot, while minimizing distance or fuel, etc.



9

Traveling Salesperson Problem

- Wait! Isn't this just a shortest path in a graph problem, like single-source shortest path, or maybe all pairs shortest paths?
- No! The TSP is much harder.
 - We're not simply looking for shortest paths.
 - We're looking for the shortest cycle that visits **all** of the cities (all of the vertices of the graph).
 - This is very different than those other shortest path problems.
 - Algorithms like Dijkstra's for single-source shortest path, or Floyd-Warshall for all-pairs shortest paths, won't help us here.
 - A* also won't help us here.

10

TSP Tour Examples



11

Informal Problem Characteristics

- Some structure to optimize.
- Value function to either minimize or maximize.
- Searching all possible configurations of the structure is intractable (e.g., TSP with n cities has $n!$ possible configurations)
- DFS is too expensive (e.g., runtime for DFS for TSP is $O(n!)$).
- No known algorithm for efficiently finding optimal.
 - Best known exact algorithms have exponential runtime or worse.
 - For example, the problem is NP Hard.
- Similar solutions have similar costs.
 - For example, if 2 tours for a TSP have $n-2$ edges in common and differ only by 2 edges, the cost of the tours will be similar.

12

Lesson 2

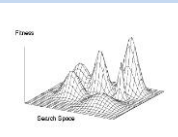
LOCAL SEARCH AND HILL CLIMBING

STOCKTON UNIVERSITY
www.stockton.edu

13

Fitness Landscape

- Fitness Landscape:
 - Imagine all possible configurations located on the surface of a landscape. The altitude of a point on the landscape corresponds to the corresponding configuration's "fitness" (i.e., value of the optimization function).
 - Similar configurations near each other on landscape.
 - Two TSP tours that differ by only 2 edges would be next to each other on landscape.
 - Term originated in biology within context of genetics.



STOCKTON UNIVERSITY
www.stockton.edu

14

Local Search

- Local Search is sometimes called Iterative Improvement.
- Basic structure of a Local Search:
 1. Start at a random initial configuration.
 2. While termination criteria not met:
 - a) Select a "neighbor".
 - b) Decide whether or not to accept that neighbor.
 3. (Some local search algorithms) Restart at step 1.
- Requires a "neighborhood function" (sometimes called "moveset operator").
- We'll look at several algorithms that fit this pattern.

STOCKTON UNIVERSITY
www.stockton.edu

15

Steepest Ascent/Descent Hill Climbing

```

X := random initial configuration
Terminate := false
While Not Terminate Do
  Terminate := true
  BestNeighbor := X
  For each neighbor X' of X do
    If (V(X') < V(BestNeighbor)) then
      Terminate = false
      BestNeighbor := X'
  If (Not Terminate) then
    X := BestNeighbor
Return X
  
```

Pseudocode written assuming minimization.

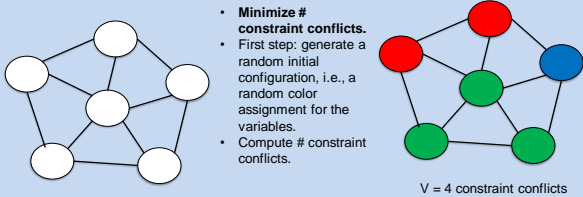
For maximization, change This < to >.

STOCKTON UNIVERSITY
www.stockton.edu

16

Example: Graph Coloring

- Minimize # constraint conflicts.
- First step: generate a random initial configuration, i.e., a random color assignment for the variables.
- Compute # constraint conflicts.



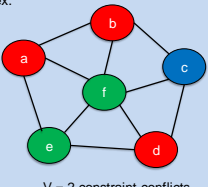
V = 4 constraint conflicts

STOCKTON UNIVERSITY
www.stockton.edu

17

Example: Graph Coloring

- Neighbor function: Change color of 1 vertex.
- Iterate over neighbors:
 - a to green, V=5
 - a to blue, V=3
 - b to green, V=4
 - b to blue, V=4
 - c to red, V=5
 - c to green, V=6
 - **d to red, V=2**
 - d to blue, V=3
 - e to red, V=3
 - e to blue, V=2
 - f to red, V=4
 - f to blue, V=3



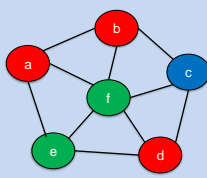
V = 2 constraint conflicts

STOCKTON UNIVERSITY
www.stockton.edu

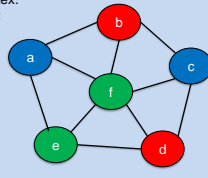
18

Example: Graph Coloring

- Neighbor function:
Change color of 1 vertex.
- Iterate over neighbors:
 - a to green, V=3
 - a to blue, V=1**
 - b to green, V=2
 - b to blue, V=2
 - c to green, V=3
 - d to green, V=4
 - d to blue, V=3
 - e to red, V=3
 - e to blue, V=1
 - f to red, V=4
 - f to blue, V=2



V = 2 constraint conflicts



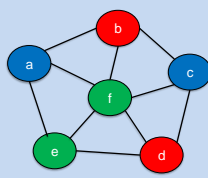
V = 1 constraint conflicts

STOCKTON UNIVERSITY
www.stockton.edu

19

Example: Graph Coloring

- Neighbor function:
Change color of 1 vertex.
- Iterate over neighbors:
 - a to red, V=2
 - a to green, V=3
 - b to green, V=2
 - b to blue, V=3
 - c to red, V=3
 - c to green, V=2
 - d to green, V=3
 - d to blue, V=2
 - e to red, V=1
 - e to blue, V=1
 - f to red, V=2
 - f to blue, V=2



V = 1 constraint conflicts

We are at a local optima so search is finished.

None of the neighbors are better than the current solution.

STOCKTON UNIVERSITY
www.stockton.edu

20

First Ascent/Descent Hill Climbing

```

X := random initial configuration
Terminate := false
While Not Terminate Do
  Terminate := true
  For each neighbor X' of X do
    If (V(X') < V(X)) then
      Terminate = false
      X := X'
      Break out of inner loop
Return X
  
```

Pseudocode written assuming minimization.

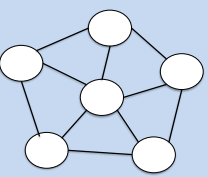
For maximization, change This < to >.

STOCKTON UNIVERSITY
www.stockton.edu

21

Example: Graph Coloring

- Minimize # constraint conflicts.
- First step: generate a random initial configuration, i.e., a random color assignment for the variables.
- Compute # constraint conflicts.



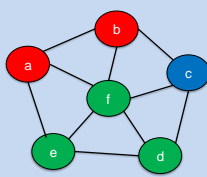
V = 4 constraint conflicts

STOCKTON UNIVERSITY
www.stockton.edu

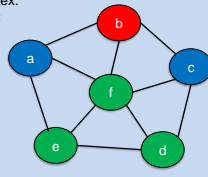
22

Example: Graph Coloring

- Neighbor function:
Change color of 1 vertex.
- Iterate over neighbors:
 - a to green, V=5
 - a to blue, V=3**



V = 4 constraint conflicts



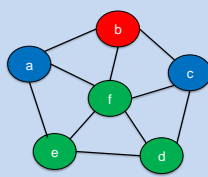
V = 3 constraint conflicts

STOCKTON UNIVERSITY
www.stockton.edu

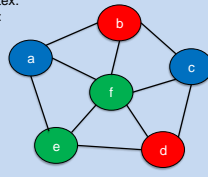
23

Example: Graph Coloring

- Neighbor function:
Change color of 1 vertex.
- Iterate over neighbors:
 - a to red, V=4
 - a to green, V=5
 - b to green, V=4
 - b to blue, V=5
 - c to red, V=4
 - c to green, V=5
 - d to red, V=1**



V = 3 constraint conflicts



V = 1 constraint conflicts

STOCKTON UNIVERSITY
www.stockton.edu

24

Example: Graph Coloring

V = 1 constraint conflicts

- Neighbor function:
 - Change color of 1 vertex.
- Iterate over neighbors:
 - a to red, V=2
 - a to green, V=3
 - b to green, V=2
 - b to blue, V=3
 - c to red, V=3
 - c to green, V=2
 - d to green, V=3
 - d to blue, V=2
 - e to red, V=1
 - e to blue, V=1
 - f to red, V=2
 - f to blue, V=2

We are at a local optima so search is finished.

None of the neighbors are better than the current solution.

STOCKTON UNIVERSITY
www.stockton.edu

25

Randomized Hill Climbing

```

X := random initial configuration
While Termination Criteria Not Met do
  X' := randomly selected neighbor of X
  If (V(X') < V(X)) then
    X := X'
Return X
  
```

Pseudocode written assuming minimization.

For maximization, change This < to >.

- When do you terminate? Here are a couple options:
 - Some maximum number of iterations of the loop.
 - Or perhaps some maximum number of iterations without accepting a neighbor.

STOCKTON UNIVERSITY
www.stockton.edu

26

Restarting a Hill Climber

- Local optima:** A solution that is at least as good as all of its neighbors.
 - Locally, you cannot do any better.
- Global optima:** A solution that is at least as good as all other solutions to the problem.
- Hill climbers find local optima, and you have no way of knowing whether or not it is the global optima.
- Common practice: Restart the hill climber multiple times.
 - Each restart finds a local optima.
 - Return the best of those local optima.

```

RestartHillClimber(P: the problem)
  X := HillClimber(P)
  for i = 1 to MAX_RESTARTS do
    X' := HillClimber(P)
    if (V(X') < V(X)) then
      X := X'
  return X
  
```

Where HillClimber refers to steepest descent, or first descent, or randomized, or any other form of hill climbing.

STOCKTON UNIVERSITY
www.stockton.edu

27

Hill Climbing Issues

- Memory: O(1)
 - Only constant memory since you only ever have two configurations in memory at any time (the current solution and one neighbor).
- Neighborhood Function
 - Critical part of application of algorithm to a problem
 - Need to determine what a neighbor is.
 - Too small a neighborhood → easy to get stuck in local optima
 - Too large a neighborhood → inefficient (loop over all neighbors)
- Tends to get stuck in local optima
 - Common Strategy: Restart when local optima reached

STOCKTON UNIVERSITY
www.stockton.edu

28

Lesson 3

LOCAL SEARCH FOR BOOLEAN SATISFIABILITY

STOCKTON UNIVERSITY
www.stockton.edu

29

Boolean Satisfiability as an Optimization Problem

- Given a Boolean expression in 3-CNF format, find an assignment of Boolean values (true/false) for the variables to maximize the number of satisfied clauses.
- Conjunctive Normal Form (CNF): Boolean expression is a conjunction (AND) of clauses, where each clause is a disjunction (OR) of Boolean literals (a variable or its negation).

$$\begin{aligned}
 &(A \vee B \vee C) \wedge (B \vee C \vee \neg D) \wedge (D \vee E \vee F) \wedge (\neg D \vee \neg E \vee \neg F) \wedge \\
 &(\neg A \vee \neg C \vee \neg F) \wedge (E \vee F \vee \neg G) \wedge (\neg A \vee G \vee \neg H) \wedge \\
 &(D \vee \neg E \vee H) \wedge (B \vee I \vee \neg J) \wedge (\neg H \vee I \vee J) \wedge (G \vee \neg I \vee \neg J)
 \end{aligned}$$

STOCKTON UNIVERSITY
www.stockton.edu

30

GSAT

- GSAT is a classic hill climbing example for Boolean Satisfiability.
 - Steepest Ascent Hill Climber
 - Neighborhood function: flip the value of any one variable
 - Flip simply means change from true to false (or false to true).
 - The G in GSAT means Greedy.

31

Walksat

- Walksat is another classic local search for Boolean Satisfiability.
- It isn't really a hill-climber, in that it sometimes chooses a neighbor that is worse than the current solution (on purpose).
- Neighborhood function:
 - Pick a random unsatisfied clause. This clause has 3 variables since we are assuming 3-SAT.
 - Consider 3 neighbors, flipping the value of exactly one of those 3 variables.
- Logic for determining which neighbor to choose:
 - If any of the 3 neighbors is better than current solution, greedily pick the best of the 3 (just like a steepest ascent hill-climber).
 - Generate a random number $r \in [0.0, 1.0)$. If $r < 0.5$, choose the least bad neighbor, otherwise choose a random neighbor.

32

Stepping through GSAT and Walksat

- We'll now step through examples of GSAT and Walksat.

33


Lesson 4

LOCAL SEARCH FOR TRAVELING SALESPERSON PROBLEM (TSP)

34


Hill Climbing for the TSP

- The Traveling Salesperson Problem
 - Find the tour (cycle which visits each exactly once) of a set of N cities that minimizes total distance.
- Configuration:
 - Permutation of the N cities
- Neighborhood function:
 - 2-change, ..., k-change



35

TSP 3-changes



36

Classic Example for TSP: k-Opt

- kOpt is a Steepest Descent Hill Climber
- Neighborhood function (for some k):
 - The set of all 2-changes, 3-changes, ..., k-changes.
 - 2-Opt neighborhood is all 2-changes, for 3-Opt, it is all 2-changes and 3-changes.
- Some (classic) Findings (Lin):
 - 3-Opt Solutions are generally much better than 2-Opt solutions.
 - 4-Opt solutions, though better than 3-Opt, are not sufficiently better to justify extra computational cost.
 - For a particular class of TSP, Lin showed that the probability that 3-Opt finds the optimal solution is: $0.5^{n/10}$ where n is number of cities.
 - Can restart R times and estimate probability of having found optimal solution: $1 - (1 - 0.5^{n/10})^R$

37

Lesson 5

ESCAPING FROM LOCAL OPTIMA

38

Local Optima

- Local Optima (definition):
 - A configuration that is at least as good as all immediate neighbors.
- Global Optima (definition):
 - A configuration that is at least as good as all possible configurations.
- Hill climbing continues until stuck:
 - We may or may not be at a global optima
 - We have no way of knowing
- Approaches to dealing with local optima:
 - Restarting (see earlier)
 - Accepting worsening neighbors sometimes (e.g., Walksat)

39

Simple Attempt at Dealing with Local Optima

```

X := random initial configuration
BestFound := X
While Termination Criteria Not Met do
    X' := randomly selected neighbor of X
    If (V(X') < V(X)) then
        X := X'
        if (V(X') < V(BestFound)) then
            BestFound := X'
    Else
        r := random real in [0.0, 1.0)
        If r < P then
            X := X'
Return BestFound
    
```

Pseudocode written
assuming minimization.

For maximization, change
This < to >.

Basic Idea: By allowing some bad
moves sometimes, might wander away
from the local optima into a more
promising area.

40

How should the probability P be set?

- Perhaps some constant like 0.1, or 0.5, etc?
- Perhaps probability that decreases over time?
 - Start $P = 1$ (i.e., accept all moves even bad ones)
 - Decrease it during search (i.e., the longer the search the less worsening moves are accepted) until $P = 0$.
 - When $P=0$, becomes a randomized hill climber by end of search.
- Perhaps probability that decreases with the "badness" of the neighbor?
 - The worse the move, the lower P is (lower probability of accepting)
 - The less bad the move, the higher P is (higher probability of accepting)
- Perhaps a combination of the above?

41

Simulated Annealing

- If random neighbor is an improvement (or of same quality), then definitely accept it.
- If random neighbor is worse than current, then accept it with the following probability:

$$e^{(V(x)-V(x'))/T_i}$$
- Known as the Boltzmann distribution
- Expressed above if minimizing, if maximizing, we use:

$$e^{(V(x')-V(x))/T_i}$$
- T is a "temperature" parameter that is "cooled" over time
- Common "cooling schedule", for cooling rate, $0 < a < 1$, and T_0 some large initial value:

$$T_i = a * T_{i-1}$$
- High temp (accept all moves), Low temp (randomized hill climbing)
- Terminate after pre-specified number of steps

42

Simulated Annealing

```

X := random initial configuration
BestFound := X
T := T0
for i = 1 to MAX_ITERATIONS do
  X' := randomly selected neighbor of X
  If (V(X') <= V(X)) then
    X := X'
    if (V(X') < V(BestFound)) then
      BestFound := X'
  Else
    r := random real in [0.0, 1.0)
    If r < e(V(x)-V(x'))/T then
      X := X'
  T := a * T
Return BestFound

```

Pseudocode written assuming minimization.

Parameters: a and T₀ are parameters, where 0.0 < a < 1.0.

a is usually 0.9 < a < 0.999

T₀ should be "high" but what this means depends on the range of the V function.

STOCKTON
UNIVERSITY
www.stockton.edu

43

More on Simulated Annealing

- Simulated Annealing introduced in 1953 by Metropolis.
 - Yes, really that long ago.
- Based on analogy to how alloys manage to find globally minimal energy level if cooled slowly.
- Much better (empirically) than hill climbing at avoiding local optima.
- Weird, provable, but not entirely useful, fact:
 - With an infinitely slow cooling rate, you'll find the global optima.
- More elaborate cooling schedules exist, including one that is "optimal" given a pre-determined search length
 - E.g., Modified Lam Schedule (Justin Boyan)
 - E.g., Lam Schedule (Lam and Delosme)

STOCKTON
UNIVERSITY
www.stockton.edu

44

Tabu Search

- Among the definitions of Tabu is "forbidden."
- Tabu Search uses a "tabu" list of configurations to try to guide the search into considering things it hasn't seen before.
- "Tabu List": Tabu search never accepts a neighbor if it is in the tabu list (unless all neighbors are Tabu).
- Tabu List has a finite, predetermined max length.
- When a neighbor is accepted, it's added to the Tabu List (removing the oldest thing from the Tabu List if it is full).
- The Tabu List is sort of like a Queue (first in first out).
- Note: There are more elaborate versions of Tabu Search.

STOCKTON
UNIVERSITY
www.stockton.edu

45

Tabu Search

```

x := random initial configuration
BestFound := x
T := a Queue initially containing only x
for i = 1 to MAX_ITERATIONS do
  If (Neighbors(x) - T ≠ ∅) then
    x := argminx' ∈ Neighbors(x) - T V(x')
    If V(x) < V(BestFound)
      BestFound := x
  Else
    x := random neighbor of x
  Add x to tail of T
  if length(T) > limit then
    remove head of T
Return BestFound

```

Pseudocode written assuming minimization.

STOCKTON
UNIVERSITY
www.stockton.edu

46